


# Binary Search Trees

A **Binary Search Tree (BST)** is a binary tree where each node contains a key such that:

$$\text{left subtree keys} < \text{node key} < \text{right subtree keys}$$

This ordering enables efficient searching, insertion, and deletion operations.

## 1 Example BST Structure



This BST contains keys ordered such that for every node:

$$\text{left subtree keys} < \text{node key} < \text{right subtree keys}$$

## 2 Use Cases

- Symbol tables in compilers
- Maintaining sorted datasets
- Searching and dynamic set operations

## 3 Time Complexity

Operation	Average Case	Worst Case (Unbalanced)
Search	$\mathcal{O}(\log n)$	$\mathcal{O}(n)$
Insert	$\mathcal{O}(\log n)$	$\mathcal{O}(n)$
Delete	$\mathcal{O}(\log n)$	$\mathcal{O}(n)$

### 3.1 Terminology

- **BST:** A binary tree with ordered keys
- **Leaf:** A node with no children

- **Internal node:** Has at least one child
- **Height:** Number of edges on the longest downward path
- **Balanced Tree:** Height is approximately  $\log n$

## 4 Why BST Operations Are $\mathcal{O}(\log n)$

In a **balanced** BST, each level of the tree halves the search space. With  $n$  nodes, the maximum height is approximately  $\log_2 n$ .

- At each comparison, we eliminate half the remaining nodes.
- This is analogous to binary search on a sorted array.

Hence, insertion, search, and deletion all take  $\mathcal{O}(\log n)$  time on average for balanced trees.

### 4.1 C++ Implementation of a BST

```
#include <iostream>
using namespace std;

struct BSTNode {
    int key;
    BSTNode* left;
    BSTNode* right;

    BSTNode(int val) : key(val), left(nullptr), right(nullptr) {}
};

class BST {
public:
    BSTNode* root;

    BST() : root(nullptr) {}

    BSTNode* insert(BSTNode* node, int key) {
        if (!node) return new BSTNode(key);
        if (key < node->key)
            node->left = insert(node->left, key);
        else if (key > node->key)
            node->right = insert(node->right, key);
        return node;
    }

    bool search(BSTNode* node, int key) const {
        if (!node) return false;
        if (key == node->key) return true;
        if (key < node->key)
            return search(node->left, key);
        else
            return search(node->right, key);
    }

    void insert(int key) {
        root = insert(root, key);
    }

    bool search(int key) const {
        return search(root, key);
    }
};
```

## 4.2 Tree Traversals

Traversal orders commonly used with BSTs:

**Inorder Traversal (Left → Root → Right):** Produces sorted order

```
void inorder(BSTNode* node) {
    if (!node) return;
    inorder(node->left);
    cout << node->key << " ";
    inorder(node->right);
}
```

**Preorder Traversal (Root → Left → Right):**

```
void preorder(BSTNode* node) {
    if (!node) return;
    cout << node->key << " ";
    preorder(node->left);
    preorder(node->right);
}
```

**Postorder Traversal (Left → Right → Root):**

```
void postorder(BSTNode* node) {
    if (!node) return;
    postorder(node->left);
    postorder(node->right);
    cout << node->key << " ";
}
```

## 5 Deletion in BSTs

To delete a node, we consider three cases:

- **Case 1: Leaf Node (0 children)** Simply remove the node.
- **Case 2: One Child** Replace the node with its only child.
- **Case 3: Two Children** Replace the node's value with either:
  - The **inorder successor** (smallest value in right subtree)
  - The **inorder predecessor** (largest value in left subtree)


Then delete the successor or predecessor node.

### Finding the Inorder Successor

To find the inorder successor of a node in a BST, we consider two cases:

- If the node has a right subtree: successor is the leftmost node in the right subtree
- If the node has no right subtree: successor is the lowest ancestor for which the node lies in its left subtree

## 5.1 Visualizing Inorder Successor Cases



**Case 1:** Node 6 has a right child (7). Its inorder successor is the leftmost node in its right subtree: 7.

**Case 2:** Node 7 has no right subtree. To find its inorder successor:

- We move up the tree from node 7 to its parent (node 6), but since 7 is a **right child**, we continue upward.
- We then move from 6 to its parent (node 3), where 6 is a **right child** again, so we continue.
- We then move from 3 to node 8, where 3 is a **left child**.
- Hence, node 8 is the lowest ancestor where node 7 (through its lineage) lies in the left subtree.

Therefore, the inorder successor of node 7 is 8.

### C++ Implementation:

#### Deletion Using inorderSuccessor

```
// Full-tree inorderSuccessor version used here
BSTNode* deleteNode(BSTNode* root, int key) {
    if (!root) return nullptr;

    if (key < root->key)
        root->left = deleteNode(root->left, key);
    else if (key > root->key)
        root->right = deleteNode(root->right, key);
    else {
        // Case 1 and 2: Node with 0 or 1 child
        if (!root->left) {
            BSTNode* temp = root->right;
            delete root;
            return temp;
        } else if (!root->right) {
            BSTNode* temp = root->left;
            delete root;
            return temp;
        }

        // Case 3: Node with 2 children
        BSTNode* succ = inorderSuccessor(root, root); // full-tree root and current node
        root->key = succ->key;
        root->right = deleteNode(root->right, succ->key);
    }
    return root;
}
```

## Deletion Function in C++

```
BSTNode* deleteNode(BSTNode* root, int key) {
    if (!root) return nullptr;

    if (key < root->key)
        root->left = deleteNode(root->left, key);
    else if (key > root->key)
        root->right = deleteNode(root->right, key);
    else {
        // Case 1 and 2: Node with 0 or 1 child
        if (!root->left) {
            BSTNode* temp = root->right;
            delete root;
            return temp;
        }
        else if (!root->right) {
            BSTNode* temp = root->left;
            delete root;
            return temp;
        }

        // Case 3: Node with 2 children
        BSTNode* succ = findMin(root->right);
        root->key = succ->key;
        root->right = deleteNode(root->right, succ->key);
    }
    return root;
}
```

**Note:** The deletion algorithm maintains the BST invariant by rearranging nodes appropriately. If balance is not preserved, performance can degrade to  $\mathcal{O}(n)$ .

## 5.2 Degenerate Cases and Balancing

A poorly structured BST can degrade to a linked list, resulting in  $\mathcal{O}(n)$  time operations. To prevent this:

- Use self-balancing trees like AVL or Red-Black Trees
- Randomize insertions or rebalance periodically

Balanced BSTs ensure height  $\mathcal{O}(\log n)$ , which guarantees logarithmic performance for search-related operations.


## 6 Tree Rotations

Tree rotations are fundamental operations used in self-balancing binary search trees (e.g., AVL and Red-Black Trees) to restore height balance after insertions or deletions. A rotation is a local restructuring operation that preserves the BST ordering.


### 6.1 Right Rotation (rotateRight)

Right rotation is applied at a node  $y$  with a left child  $x$  to reduce the height of the left subtree.

**Before:**




**After rotateRight(y):**




## 6.2 Left Rotation (rotateLeft)

Left rotation is applied at a node  $x$  with a right child  $y$  to reduce the height of the right subtree.

**Before:**



**After rotateLeft(x):**



## 6.3 Preserving the BST Property During Rotations

Tree rotations maintain the binary search tree property:

$$\text{left subtree} < \text{node} < \text{right subtree}$$


Consider the right rotation at node  $y$  (where  $x$  is its left child):

**Before rotation:**

- Subtree A: All keys  $< x.\text{key}$

- Subtree B:  $x.\text{key} < \text{keys} < y.\text{key}$
- Subtree C: All keys  $> y.\text{key}$

**Right Rotation:**



**After rotation:**

- Subtree A remains in the left of x
- Subtree B becomes left of y, and since B's keys were  $> x.\text{key}$  and  $< y.\text{key}$ , they remain valid
- Subtree C remains to the right of y

All relative key orderings are preserved, satisfying:

$$A < x < B < y < C$$

Both left and right rotations preserve in-order traversal and relative ordering of keys. Hence, the BST property remains intact after any single rotation.

## 7 Practice problem

Leetcode Problem 108: *Convert Sorted Array to Binary Search Tree*

<https://leetcode.com/problems/convert-sorted-array-to-binary-search-tree/description/>