

# Graph Algorithms

## 1 What is a Graph?

A **graph** is a mathematical structure used to model pairwise relations between objects. A graph  $G$  consists of:

- A set of **vertices** (or nodes),  $V$ .
- A set of **edges**,  $E$ , representing connections between pairs of vertices.

Graphs may be:


- **Undirected** or **Directed**
- **Weighted** or **Unweighted**
- **Cyclic** or **Acyclic**
- **Connected** or **Disconnected**

## 2 Why Graphs?

Graphs are widely used in:

- Computer networks (routers as vertices, links as edges)
- Social networks (people as vertices, relationships as edges)
- Maps and GPS routing (cities and roads)
- Scheduling and dependency analysis

## 3 Example Graph



## 4 Common Terminology

- **Vertex (Node):** A point in the graph.
- **Edge:** A connection between two vertices.
- **Degree:** Number of edges incident to a vertex.

- **Path:** A sequence of edges connecting a sequence of vertices.
- **Cycle:** A path that begins and ends at the same vertex.
- **Connected Graph:** A graph where there's a path between any two vertices.


## 5 Directed vs. Undirected Graphs

- In an **undirected graph**, edges have no direction. An edge  $(u, v)$  implies that  $u$  is connected to  $v$  and vice versa.
- In a **directed graph (digraph)**, each edge has a direction. An edge  $(u, v)$  goes from  $u$  to  $v$  only.

### Edge Weights

Some graphs associate a numerical value (called a **weight**) with each edge. These are known as **weighted graphs**. Weights might represent distances, costs, capacities, etc.

### Illustration



#### Notes:

- The arrows indicate that the graph is directed.
- Numbers on the edges represent weights.
- If the same edges had no arrows and were bidirectional, it would represent an undirected weighted graph.

## 6 Graph Representation

### 1. Adjacency Matrix

- A 2D matrix of size  $n \times n$ , where  $n$  is the number of vertices.
- Entry  $(i, j)$  is 1 (or weight  $w$ ) if there is an edge from  $i$  to  $j$ .

```
// Adjacency Matrix Representation
int n = 4; // number of vertices
vector<vector<int>> adjMatrix(n, vector<int>(n, 0));

// Add edge between u and v
adjMatrix[u][v] = 1;
adjMatrix[v][u] = 1; // if undirected
```

### 2. Adjacency List

- An array of lists, where index  $i$  stores the list of adjacent vertices to  $i$ .
- More space-efficient for sparse graphs.

```
// Adjacency List Representation
int n = 4; // number of vertices
vector<vector<int>> adjList(n);

// Add edge between u and v
adjList[u].push_back(v);
adjList[v].push_back(u); // if undirected
```

## 7 Breadth-First Search (BFS)

**BFS** explores a graph level by level, using a queue. It is useful for finding the shortest path in unweighted graphs.

```
// BFS from source vertex
void bfs(int src, vector<vector<int>>& adj, int n) {
    vector<bool> visited(n, false);
    queue<int> q;

    visited[src] = true;
    q.push(src);

    while (!q.empty()) {
        int node = q.front(); q.pop();
        cout << node << " ";

        for (int neighbor : adj[node]) {
            if (!visited[neighbor]) {
                visited[neighbor] = true;
                q.push(neighbor);
            }
        }
    }
}
```

**Time Complexity:**  $O(V + E)$

## 8 Depth-First Search (DFS)

**DFS** explores as far as possible along a branch before backtracking. It can be implemented recursively or using a stack.

```
// DFS using recursion
void dfs(int node, vector<vector<int>>& adj, vector<bool>& visited) {
    visited[node] = true;
    cout << node << " ";

    for (int neighbor : adj[node]) {
        if (!visited[neighbor]) {
            dfs(neighbor, adj, visited);
        }
    }
}
```

**Time Complexity:**  $O(V + E)$

## 9 Applications of BFS and DFS

### 9.1 BFS for Unweighted Shortest Path

Breadth-First Search (BFS) can be used to compute the shortest path from a source vertex to all other vertices in an **unweighted graph**.

The idea is that BFS explores vertices in layers: it visits all nodes at distance 1 before distance 2, and so on. Thus, the first time we visit a node, we've found the shortest path to it in terms of number of edges.

```
// Compute shortest path from source using BFS
vector<int> bfsShortestPath(int src, const vector<vector<int>>& adj, int n) {
    vector<int> dist(n, -1); // -1 indicates unreachable
    queue<int> q;

    dist[src] = 0;
    q.push(src);

    while (!q.empty()) {
        int u = q.front(); q.pop();
        for (int v : adj[u]) {
            if (dist[v] == -1) {
                dist[v] = dist[u] + 1;
                q.push(v);
            }
        }
    }
    return dist;
}
```

**Note:** This approach does not handle edge weights. Use Dijkstra's algorithm if weights are involved.

### 9.2 Checking Graph Connectivity

To determine if a graph is **connected**, we can use either BFS or DFS:

- Start from any node and perform BFS or DFS.
- After traversal, check if all nodes were visited.

```
// Simple DFS function for connectivity
void dfs(int u, const vector<vector<int>>& adj, vector<bool>& visited, int& count) {
    visited[u] = true;
    count++;
    for (int v : adj[u]) {
        if (!visited[v]) {
            dfs(v, adj, visited, count);
        }
    }
}

// Returns true if the graph is connected
bool isConnected(const vector<vector<int>>& adj, int n) {
    vector<bool> visited(n, false);
    int count = 0;

    dfs(0, adj, visited, count); // Start from node 0

    return count == n;
}
```

**Note on Directed Graphs:**

- A directed graph is **strongly connected** if there is a path from every vertex to every other vertex following the direction of edges.

- It is **weakly connected** if the graph is connected when edge directions are ignored (i.e., treated as undirected).

Determining strong connectivity requires more advanced algorithms such as Kosaraju's or Tarjan's algorithm.

### 9.3 Topological Sorting via DFS

Topological sorting is applicable to **Directed Acyclic Graphs (DAGs)**. It produces a linear ordering of vertices such that for every directed edge  $(u \rightarrow v)$ , vertex  $u$  appears before  $v$  in the ordering.

DFS-based topological sorting:

- Perform DFS on the graph.
- On finishing a node, push it to a stack or prepend to a list.
- Reverse the result at the end.

```
// Topological sort using DFS
void topoDFS(int u, vector<vector<int>>& adj, vector<bool>& visited, vector<int>& result) {
    visited[u] = true;
    for (int v : adj[u]) {
        if (!visited[v])
            topoDFS(v, adj, visited, result);
    }
    result.push_back(u); // Add after visiting children
}

vector<int> topologicalSort(int n, vector<vector<int>>& adj) {
    vector<bool> visited(n, false);
    vector<int> result;

    for (int i = 0; i < n; ++i) {
        if (!visited[i])
            topoDFS(i, adj, visited, result);
    }

    reverse(result.begin(), result.end());
    return result;
}
```

### 9.4 Spanning Trees and BFS/DFS Construction

A **spanning tree** of a connected, undirected graph is a subgraph that:

- Includes all the vertices of the original graph,
- Is a tree (i.e., contains no cycles),
- Has exactly  $n - 1$  edges if there are  $n$  vertices.

Spanning trees are fundamental in graph theory and network design (e.g., building minimum-cost networks without cycles).

**Key Property:** Every connected undirected graph has at least one spanning tree.

#### Constructing a Spanning Tree Using BFS

You can build a spanning tree by performing a BFS and recording the edges used to first discover each vertex.

```
// Returns list of edges in the BFS spanning tree
vector<pair<int, int>> bfsSpanningTree(int start, const vector<vector<int>>& adj, int n) {
    vector<bool> visited(n, false);
    vector<pair<int, int>> treeEdges;
    queue<int> q;

    visited[start] = true;
    q.push(start);

    while (!q.empty()) {
        int u = q.front(); q.pop();
        for (int v : adj[u]) {
            if (!visited[v]) {
                visited[v] = true;
                treeEdges.push_back({u, v});
                q.push(v);
            }
        }
    }

    return treeEdges;
}
```

## Constructing a Spanning Tree Using DFS

Likewise, DFS can be used to form a spanning tree by recording edges used during the traversal.

```
// Returns list of edges in the DFS spanning tree
void dfsSpanningTree(int u, const vector<vector<int>>& adj, vector<bool>& visited,
    vector<pair<int, int>>& treeEdges) {
    visited[u] = true;
    for (int v : adj[u]) {
        if (!visited[v]) {
            treeEdges.push_back({u, v});
            dfsSpanningTree(v, adj, visited, treeEdges);
        }
    }
}
```

To use:

```
vector<bool> visited(n, false);
vector<pair<int, int>> treeEdges;
dfsSpanningTree(0, adj, visited, treeEdges);
```

**Note:** These methods generate *some* spanning tree, not necessarily a minimum spanning tree (MST). For MSTs, use Prim's or Kruskal's algorithm.

## 10 Advanced Graph Algorithms: Greedy Approach

Many important graph problems can be solved efficiently using the **greedy paradigm**, where a local optimal choice is made at each step with the hope that it leads to a global optimum.

Two key graph algorithms based on this idea are:

- **Prim's Algorithm** for finding a Minimum Spanning Tree (MST),
- **Dijkstra's Algorithm** for finding the shortest path from a source node.

### 10.1 Prim's Algorithm: Minimum Spanning Tree

Prim's algorithm constructs a **minimum spanning tree** by growing the tree one edge at a time. At each step, it adds the minimum weight edge that connects a vertex in the tree to a vertex outside the tree.

**Greedy Principle:** At each step, choose the lightest edge that expands the tree without forming a cycle.

```
// Prim's algorithm without priority queue
vector<int> primMST(const vector<vector<pair<int, int>>>& adj, int n) {
    vector<bool> inMST(n, false);
    vector<int> key(n, INT_MAX); // Minimum weight to include node
    vector<int> parent(n, -1);   // Store MST edges
    key[0] = 0;

    for (int count = 0; count < n; ++count) {
        // Find the vertex u not in MST with minimum key[u]
        int u = -1;
        for (int i = 0; i < n; ++i) {
            if (!inMST[i] && (u == -1 || key[i] < key[u]))
                u = i;
        }

        inMST[u] = true;

        for (auto [v, weight] : adj[u]) {
            if (!inMST[v] && weight < key[v]) {
                key[v] = weight;
                parent[v] = u;
            }
        }
    }

    return parent; // parent[i] gives the MST edge to i
}
```

**Time Complexity:**  $O(V^2)$  due to the repeated minimum selection via linear scan.

**Time Complexity:**  $O((V + E) \log V)$  with a priority queue. **Input:** Undirected, connected, weighted graph.

## 10.2 Dijkstra's Algorithm: Single-Source Shortest Path

Dijkstra's algorithm finds the shortest path from a source vertex to all other vertices in a graph with non-negative edge weights.

**Greedy Principle:** At each step, pick the closest unvisited vertex and update distances to its neighbors.

```
// Dijkstra's algorithm without priority queue
vector<int> dijkstra(int src, const vector<vector<pair<int, int>>>& adj, int n) {
    vector<int> dist(n, INT_MAX);
    vector<bool> visited(n, false);
    dist[src] = 0;

    for (int count = 0; count < n; ++count) {
        // Find unvisited vertex u with smallest dist[u]
        int u = -1;
        for (int i = 0; i < n; ++i) {
            if (!visited[i] && (u == -1 || dist[i] < dist[u]))
                u = i;
        }

        visited[u] = true;

        for (auto [v, weight] : adj[u]) {
            if (dist[u] + weight < dist[v]) {
                dist[v] = dist[u] + weight;
            }
        }
    }

    return dist;
}
```

**Time Complexity:**  $O(V^2)$  — acceptable for dense graphs or small inputs.

### 10.3 Why Greedy Works

Both Prim's and Dijkstra's algorithms build solutions incrementally:

- They always pick the next node or edge that seems best *at the moment*.
- They avoid revisiting or backtracking once a decision is made.
- Correctness is guaranteed due to the properties of MSTs (Prim) and non-negative weights (Dijkstra).

### 10.4 How Heaps Improve Efficiency

Both Prim's and Dijkstra's algorithms involve repeatedly selecting the vertex with the smallest key/distance. This operation is costly with a linear scan ( $O(V)$  per iteration).

By using a **min-heap (priority queue)**, we reduce this selection time to  $O(\log V)$ :

- **Prim's with heap:**  $O((V + E) \log V)$
- **Dijkstra with heap:**  $O((V + E) \log V)$

## 11 Other Graph Algorithms (Brief Overview)

- **Kruskal's Algorithm:** Greedy algorithm for finding a minimum spanning tree using edge sorting and union-find.
- **Bellman-Ford Algorithm:** Computes shortest paths from a source, allowing negative edge weights.
- **Floyd-Warshall Algorithm:** Computes all-pairs shortest paths using dynamic programming.
- **Topological Sort (Kahn's Algorithm):** Iterative method to produce a topological ordering of a DAG.
- **Kosaraju's Algorithm:** Detects strongly connected components in a directed graph.
- **Tarjan's Algorithm:** Finds strongly connected components in a single DFS traversal.
- **Union-Find (Disjoint Set Union):** Efficient structure for handling connectivity queries and cycle detection.
- **Hamiltonian Cycle:** A cycle that visits each vertex exactly once and returns to the starting vertex.
- **Eulerian Cycle:** A cycle that visits every edge exactly once and returns to the start (exists if all degrees are even in an undirected graph).